

In [1]:

```

1 #Importing our required packages Packages
2 import numpy as np
3 import pandas as pd
4 import keras
5 from keras.models import Sequential
6 from keras.layers import Dense, Flatten, Conv2D, MaxPool2D, Dropout
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9 from tensorflow.keras.preprocessing.image import ImageDataGenerator
10

```

In [2]:

```

1 #Loading our datasets,
2 train_df=pd.read_csv('datasets/sign_mnist_train.csv')
3 test_df=pd.read_csv('datasets/sign_mnist_test.csv')

```

In [5]:

```


1 #Looking at the basic statistics of our dataset.
2 train_df.describe()

```

Out[5]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5
count	27455.000000	27455.000000	27455.000000	27455.000000	27455.000000	27455.000000
mean	12.318813	145.419377	148.500273	151.247714	153.546531	156.210891
std	7.287552	41.358555	39.942152	39.056286	38.595247	37.111165
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	6.000000	121.000000	126.000000	130.000000	133.000000	137.000000
50%	13.000000	150.000000	153.000000	156.000000	158.000000	160.000000
75%	19.000000	174.000000	176.000000	178.000000	179.000000	181.000000
max	24.000000	255.000000	255.000000	255.000000	255.000000	255.000000

8 rows × 785 columns



In [25]:

```
1 #taking a look at our data
2 train_df.head()
```

Out[25]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776
0	3	107	118	127	134	139	143	146	150	153	...	207	208
1	6	155	157	156	156	156	157	156	158	158	...	69	149
2	2	187	188	188	187	187	186	187	188	187	...	202	203
3	2	211	211	212	212	211	210	211	210	210	...	235	236
4	13	164	167	170	172	176	179	180	184	185	...	92	109

5 rows × 785 columns



In [9]:

```
1 train_label=train_df['label']
2 train_label.head()
3 trainset=train_df.drop(['label'],axis=1)
4 trainset.head()
```

Out[9]:

	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	...	pixel775	pixel776
0	107	118	127	134	139	143	146	150	153	156	...	207	208
1	155	157	156	156	156	157	156	158	158	157	...	69	149
2	187	188	188	187	187	186	187	188	187	186	...	202	203
3	211	211	212	212	211	210	211	210	210	211	...	235	236
4	164	167	170	172	176	179	180	184	185	186	...	92	109

5 rows × 784 columns



In [10]:

```
1 X_train = trainset.values
2 X_train = trainset.values.reshape(-1,28,28,1)
3 print(X_train.shape)
```

(27455, 28, 28, 1)

In [11]:

```

1 test_label=test_df['label']
2 X_test=test_df.drop(['label'],axis=1)
3 print(X_test.shape)
4 X_test.head()

```

(7172, 784)

Out[11]:

	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	...	pixel775	pixel776
0	149	149	150	150	150	151	151	150	151	152	...	138	139
1	126	128	131	132	133	134	135	135	136	138	...	47	48
2	85	88	92	96	105	123	135	143	147	152	...	68	69
3	203	205	207	206	207	209	210	209	210	209	...	154	155
4	188	191	193	195	199	201	202	203	203	203	...	26	27

5 rows × 784 columns

In [12]:

```

1 #Converting our integers to binary form. With the help of LabelBinarizer
2 from sklearn.preprocessing import LabelBinarizer
3 lb=LabelBinarizer()
4 y_train=lb.fit_transform(train_label)
5 y_test=lb.fit_transform(test_label)

```

In [13]:

```
1 y_train
```

Out[13]:

```

array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 1, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 1, 0]])

```

In [14]:

```
1 X_test=X_test.values.reshape(-1,28,28,1)
```

In [15]:

```
1 print(X_train.shape,y_train.shape,X_test.shape,y_test.shape)
```

(27455, 28, 28, 1) (27455, 24) (7172, 28, 28, 1) (7172, 24)

In [16]:

```
1  '''
2  Augmenting the image dataset to generate new data
3
4  ImageDataGenerator package from keras.preprocessing.image allows to add different d
5
6  Here is the package details https://keras.io/preprocessing/image/
7
8  The image dataset is also normalised here using the rescale parameter which divides
9  '''
10
11  train_datagen = ImageDataGenerator(rescale = 1./255,
12                                     rotation_range = 0,
13                                     height_shift_range=0.2,
14                                     width_shift_range=0.2,
15                                     shear_range=0,
16                                     zoom_range=0.2,
17                                     horizontal_flip=True,
18                                     fill_mode='nearest')
19
20  X_test=X_test/255
```

In [17]:


```

1  '''
2  Visualization of the Dataset
3  Preview of the images in the training dataset
4  '''
5  fig,axe=plt.subplots(2,2)
6  fig.suptitle('Preview of dataset')
7  axe[0,0].imshow(X_train[0].reshape(28,28),cmap='gray')
8  axe[0,0].set_title('label: 3  letter: C')
9  axe[0,1].imshow(X_train[1].reshape(28,28),cmap='gray')
10 axe[0,1].set_title('label: 6  letter: F')
11 axe[1,0].imshow(X_train[2].reshape(28,28),cmap='gray')
12 axe[1,0].set_title('label: 2  letter: B')
13 axe[1,1].imshow(X_train[4].reshape(28,28),cmap='gray')
14 axe[1,1].set_title('label: 13  letter: M')

```

Out[17]:

Text(0.5, 1.0, 'label: 13 letter: M')



In [19]:

```

1  '''
2  Building our CNN Model
3  This model consist of-
4
5  1. Three convolution layer which uses MaxPooling for better feature capture.
6  2. A dense layer with 512 units
7  3. The output layer consist of 24 units for 24 different classes
8
9
10
11 Some information about the Convolution layers
12 The activation fucntion which we have used is ReLu.
13
14 Conv layer 1 -- UNITS - 128 KERNEL SIZE - 5 * 5 STRIDE LENGTH - 1 ACTIVATION - ReLu
15 Conv layer 2 -- UNITS - 64 KERNEL SIZE - 3 * 3 STRIDE LENGTH - 1 ACTIVATION - ReLu
16 Conv layer 3 -- UNITS - 32 KERNEL SIZE - 2 * 2 STRIDE LENGTH - 1 ACTIVATION - ReLu
17
18 MaxPool layer 1 -- MAX POOL WINDOW - 3 * 3 STRIDE - 2
19 MaxPool layer 2 -- MAX POOL WINDOW - 2 * 2 STRIDE - 2
20 MaxPool layer 3 -- MAX POOL WINDOW - 2 * 2 STRIDE - 2
21 '''
22
23 model=Sequential()
24 model.add(Conv2D(128,kernel_size=(5,5),
25                 strides=1,padding='same',activation='relu',input_shape=(28,28,1)))
26 model.add(MaxPool2D(pool_size=(3,3),strides=2,padding='same'))
27 model.add(Conv2D(64,kernel_size=(2,2),
28                 strides=1,activation='relu',padding='same'))
29 model.add(MaxPool2D((2,2),2,padding='same'))
30 model.add(Conv2D(32,kernel_size=(2,2),
31                 strides=1,activation='relu',padding='same'))
32 model.add(MaxPool2D((2,2),2,padding='same'))
33
34 model.add(Flatten())

```

In [20]:

```

1 #Dense and output layers
2 model.add(Dense(units=512,activation='relu'))
3 model.add(Dropout(rate=0.25))
4 model.add(Dense(units=24,activation='softmax'))
5 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 128)	3328
max_pooling2d (MaxPooling2D)	(None, 14, 14, 128)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	32832
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 32)	8224
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 32)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 512)	262656
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 24)	12312
=====		
Total params: 319,352		
Trainable params: 319,352		
Non-trainable params: 0		

In [21]:

```

1 model.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])

```

In [22]:

```

1 #Training the model
2 model.fit(train_datagen.flow(X_train,y_train,batch_size=200),
3           epochs = 35,
4           validation_data=(X_test,y_test),
5           shuffle=1
6           )

```

Epoch 1/35

138/138 [=====] - 89s 620ms/step - loss: 2.992
 8 - accuracy: 0.1007 - val_loss: 2.3336 - val_accuracy: 0.2716

Epoch 2/35

138/138 [=====] - 85s 616ms/step - loss: 2.265
 2 - accuracy: 0.2861 - val_loss: 1.5806 - val_accuracy: 0.4558

Epoch 3/35

138/138 [=====] - 97s 702ms/step - loss: 1.660
 4 - accuracy: 0.4520 - val_loss: 1.1958 - val_accuracy: 0.5471

Epoch 4/35

138/138 [=====] - 111s 803ms/step - loss: 1.28
 77 - accuracy: 0.5664 - val_loss: 0.9396 - val_accuracy: 0.6528

Epoch 5/35

138/138 [=====] - 108s 786ms/step - loss: 1.06
 20 - accuracy: 0.6426 - val_loss: 0.7139 - val_accuracy: 0.7400

Epoch 6/35

138/138 [=====] - 106s 770ms/step - loss: 0.88
 63 - accuracy: 0.6988 - val_loss: 0.5081 - val_accuracy: 0.8260

Epoch 7/35

138/138 [=====] - 98s 711ms/step - loss: 0.766
 9 - accuracy: 0.7396 - val_loss: 0.4069 - val_accuracy: 0.8639

Epoch 8/35

138/138 [=====] - 103s 747ms/step - loss: 0.65
 12 - accuracy: 0.7803 - val_loss: 0.3731 - val_accuracy: 0.8681

Epoch 9/35

138/138 [=====] - 103s 750ms/step - loss: 0.57
 92 - accuracy: 0.8017 - val_loss: 0.2901 - val_accuracy: 0.9091

Epoch 10/35

138/138 [=====] - 112s 811ms/step - loss: 0.50
 62 - accuracy: 0.8271 - val_loss: 0.2422 - val_accuracy: 0.9285

Epoch 11/35

138/138 [=====] - 104s 752ms/step - loss: 0.45
 52 - accuracy: 0.8467 - val_loss: 0.1882 - val_accuracy: 0.9472

Epoch 12/35

138/138 [=====] - 111s 807ms/step - loss: 0.41
 73 - accuracy: 0.8578 - val_loss: 0.1622 - val_accuracy: 0.9505

Epoch 13/35

138/138 [=====] - 109s 792ms/step - loss: 0.38
 06 - accuracy: 0.8715 - val_loss: 0.1261 - val_accuracy: 0.9672

Epoch 14/35

138/138 [=====] - 111s 803ms/step - loss: 0.34
 53 - accuracy: 0.8836 - val_loss: 0.1193 - val_accuracy: 0.9670

Epoch 15/35

138/138 [=====] - 102s 738ms/step - loss: 0.31
 59 - accuracy: 0.8937 - val_loss: 0.0935 - val_accuracy: 0.9742

Epoch 16/35

138/138 [=====] - 107s 774ms/step - loss: 0.31
 54 - accuracy: 0.8920 - val_loss: 0.1538 - val_accuracy: 0.9543

Epoch 17/35

138/138 [=====] - 105s 759ms/step - loss: 0.28
 14 - accuracy: 0.9061 - val_loss: 0.0865 - val_accuracy: 0.9785


```
Epoch 18/35
138/138 [=====] - 111s 805ms/step - loss: 0.26
05 - accuracy: 0.9114 - val_loss: 0.0787 - val_accuracy: 0.9763
Epoch 19/35
138/138 [=====] - 112s 807ms/step - loss: 0.24
16 - accuracy: 0.9185 - val_loss: 0.0570 - val_accuracy: 0.9831
Epoch 20/35
138/138 [=====] - 109s 786ms/step - loss: 0.23
36 - accuracy: 0.9219 - val_loss: 0.0634 - val_accuracy: 0.9826
Epoch 21/35
138/138 [=====] - 78s 559ms/step - loss: 0.206
8 - accuracy: 0.9286 - val_loss: 0.0500 - val_accuracy: 0.9894
Epoch 22/35
138/138 [=====] - 76s 549ms/step - loss: 0.201
9 - accuracy: 0.9301 - val_loss: 0.0622 - val_accuracy: 0.9808
Epoch 23/35
138/138 [=====] - 76s 552ms/step - loss: 0.198
4 - accuracy: 0.9346 - val_loss: 0.0323 - val_accuracy: 0.9948
Epoch 24/35
138/138 [=====] - 78s 562ms/step - loss: 0.190
8 - accuracy: 0.9363 - val_loss: 0.0232 - val_accuracy: 0.9967
Epoch 25/35
138/138 [=====] - 76s 550ms/step - loss: 0.193
7 - accuracy: 0.9348 - val_loss: 0.0653 - val_accuracy: 0.9802
Epoch 26/35
138/138 [=====] - 76s 548ms/step - loss: 0.173
8 - accuracy: 0.9415 - val_loss: 0.0278 - val_accuracy: 0.9957
Epoch 27/35
138/138 [=====] - 76s 549ms/step - loss: 0.156
8 - accuracy: 0.9485 - val_loss: 0.0217 - val_accuracy: 0.9974
Epoch 28/35
138/138 [=====] - 76s 552ms/step - loss: 0.162
4 - accuracy: 0.9456 - val_loss: 0.0291 - val_accuracy: 0.9915
Epoch 29/35
138/138 [=====] - 77s 555ms/step - loss: 0.153
5 - accuracy: 0.9490 - val_loss: 0.0246 - val_accuracy: 0.9927
Epoch 30/35
138/138 [=====] - 76s 550ms/step - loss: 0.148
4 - accuracy: 0.9488 - val_loss: 0.0221 - val_accuracy: 0.9932
Epoch 31/35
138/138 [=====] - 76s 550ms/step - loss: 0.154
0 - accuracy: 0.9484 - val_loss: 0.0375 - val_accuracy: 0.9880
Epoch 32/35
138/138 [=====] - 76s 553ms/step - loss: 0.137
4 - accuracy: 0.9546 - val_loss: 0.0314 - val_accuracy: 0.9883
Epoch 33/35
138/138 [=====] - 77s 554ms/step - loss: 0.132
1 - accuracy: 0.9565 - val_loss: 0.0443 - val_accuracy: 0.9861
Epoch 34/35
138/138 [=====] - 76s 552ms/step - loss: 0.129
5 - accuracy: 0.9567 - val_loss: 0.0226 - val_accuracy: 0.9886
Epoch 35/35
138/138 [=====] - 76s 552ms/step - loss: 0.120
8 - accuracy: 0.9585 - val_loss: 0.0217 - val_accuracy: 0.9954
```

Out[22]:

<keras.callbacks.History at 0x21c4fb80550>

In [23]:

```
1 #Evaluating the model
2 (ls,acc)=model.evaluate(x=X_test,y=y_test)
3
```

225/225 [=====] - 6s 27ms/step - loss: 0.0217 - a
ccuracy: 0.9954

In [24]:

```
1 print('Model Accuracy is {}'.format(acc*100))
2
```

Model Accuracy is 99.5398759841919%

In []:

```
1
```